

5 Normal Multivariate Distribution

MA2003B Application of Multivariate Methods
in Data Science

Raul Gomez



Tecnológico
de Monterrey

Definition and Basic Properties

Definition

Let μ be a (row) vector and let Σ be a symmetric positive definite matrix. We say that a random (row) vector X has **normal multivariate distribution** with mean μ and covariance matrix Σ , in symbols $X \sim \mathcal{N}_p(\mu, \Sigma)$, if its probability density is given by

$$\frac{1}{(2\pi)^{p/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(X - \mu)\Sigma^{-1}(X - \mu)^T\right) dX$$

Remark

Some things to notice about the definition above:

1. Just like in the univariate case, the normal multivariate distribution is completely determined by its mean μ and covariance matrix Σ .

Remark

Some things to notice about the definition above:

1. Just like in the univariate case, the normal multivariate distribution is completely determined by its mean μ and covariance matrix Σ .
2. The covariance matrix Σ must be symmetric and positive definite, which ensures that the distribution is well-defined and that the exponent in the density function is negative definite.

Remark

Some things to notice about the definition above:

1. Just like in the univariate case, the normal multivariate distribution is completely determined by its mean μ and covariance matrix Σ .
2. The covariance matrix Σ must be symmetric and positive definite, which ensures that the distribution is well-defined and that the exponent in the density function is negative definite.
3. The term

$$|\Sigma| = \det(\Sigma)$$

is the **generalized variance** of the distribution. It can be interpreted as a measure of the "spread" of the distribution in the p -dimensional space.

Remark

Some things to notice about the definition above:

1. Just like in the univariate case, the normal multivariate distribution is completely determined by its mean μ and covariance matrix Σ .
2. The covariance matrix Σ must be symmetric and positive definite, which ensures that the distribution is well-defined and that the exponent in the density function is negative definite.
3. The term

$$|\Sigma| = \det(\Sigma)$$

is the **generalized variance** of the distribution. It can be interpreted as a measure of the "spread" of the distribution in the p -dimensional space.

4. Notice that the term $(X - \mu)\Sigma^{-1}(X - \mu)^T$ is the square of the Mahalanobis distance between the random vector X and the mean μ .

Just like in the univariate case, the normal multivariate distribution is singled out among all other possible distributions by the Central Limit Theorem, that roughly states that the sum of a large number of independent and identically distributed random variables (properly standardized) converges to a normal distribution.

Multivariate Central Limit Theorem

Theorem

Let $\{X_i\}_{i=1}^{\infty}$ be a sequence of independent and identically distributed random vectors in $\mathbb{R}^p$ with mean vector μ and covariance matrix Σ . Define the sample mean

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i.$$

Then, as $n \rightarrow \infty$,

$$\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} \mathcal{N}_p(0, \Sigma),$$

where $\xrightarrow{d}$ denotes convergence in distribution and $\mathcal{N}_p(0, \Sigma)$ is the p -dimensional multivariate normal distribution with mean vector 0 and covariance matrix Σ .

Remark

Even though the Central Limit Theorem singles out the normal multivariate distribution as a particularly important distribution, it is not the only relevant distribution in multivariate statistics.

Remark

Even though the Central Limit Theorem singles out the normal multivariate distribution as a particularly important distribution, it is not the only relevant distribution in multivariate statistics.

Actually, most real-world data does not follow a normal distribution, and many statistical methods are designed to handle non-normal data.

Remark

Even though the Central Limit Theorem singles out the normal multivariate distribution as a particularly important distribution, it is not the only relevant distribution in multivariate statistics.

Actually, most real-world data does not follow a normal distribution, and many statistical methods are designed to handle non-normal data.

Furthermore, important statistical methods, like linear regression, principal component analysis and clustering do not yield good results when applied to a normal multivariate distribution.

Normal Multivariate Contours

Notice that the probability density function of the normal multivariate distribution is a Gaussian function in p dimensions.

Normal Multivariate Contours

Notice that the probability density function of the normal multivariate distribution is a Gaussian function in p dimensions.

This means that the contours of the distribution are ellipsoids centered at the mean μ and whose shape is controlled by Σ .

Normal Multivariate Contours

Notice that the probability density function of the normal multivariate distribution is a Gaussian function in p dimensions.

This means that the contours of the distribution are ellipsoids centered at the mean μ and whose shape is controlled by Σ .

The probability distribution of these contours is controlled by the χ^2 distribution in the following sense:

Let $0 \leq \varpi \leq 1$ be a probability level, and let c be such that

$$\mathbb{P}(\chi_p^2 \leq c) = \varpi$$

Let $0 \leq \varpi \leq 1$ be a probability level, and let c be such that

$$\mathbb{P}(\chi_p^2 \leq c) = \varpi$$

Then, the contour that contains a proportion ϖ of the probability mass of the distribution is given by the set of points X that satisfy the equation

$$(X - \mu)\Sigma^{-1}(X - \mu)^T = c$$

In the following figure, we show the contours of the bivariate normal distribution with different covariance matrices.

In the following figure, we show the contours of the bivariate normal distribution with different covariance matrices.

The contours are ellipses centered at the mean $\mu = (0,0)$, and their shape is determined by the covariance matrix Σ .

In the following figure, we show the contours of the bivariate normal distribution with different covariance matrices.

The contours are ellipses centered at the mean $\mu = (0,0)$, and their shape is determined by the covariance matrix Σ .

The color gradient represents the density of the distribution, with darker colors indicating higher density.

Contours of the Bivariate Normal




Figure 1: Bivariate Normal Distribution with Different Covariance Matrices

Standardization of the Normal Multivariate Distribution

Just like in the univariate case, given a random vector $X \sim \mathcal{N}_p(\mu, \Sigma)$, we can standardize it to obtain a new random vector Z that follows the standard multivariate normal distribution $\mathcal{N}_p(0, I_p)$, where I_p is the identity matrix of size p .

Standardization of the Normal Multivariate Distribution

Just like in the univariate case, given a random vector $X \sim \mathcal{N}_p(\mu, \Sigma)$, we can standardize it to obtain a new random vector Z that follows the standard multivariate normal distribution $\mathcal{N}_p(0, I_p)$, where I_p is the identity matrix of size p .

To describe this standardization process, we first need to remember some linear algebra results.

Singular Value Decomposition

Theorem

Let Σ be a symmetric positive definite matrix of size $p \times p$. Then, there exists an orthogonal matrix Q and a diagonal matrix Λ with positive entries such that

$$\Sigma = Q\Lambda Q^T$$

Singular Value Decomposition

Theorem

Let Σ be a symmetric positive definite matrix of size $p \times p$. Then, there exists an orthogonal matrix Q and a diagonal matrix Λ with positive entries such that

$$\Sigma = Q\Lambda Q^T$$

Remark

Remember that a matrix Q is orthogonal if $Q^T Q = I_p$, where I_p is the identity matrix of size p . This means that the columns of Q are orthonormal vectors, and thus they form a basis for $\mathbb{R}^p$.

Remark

The singular value decomposition (SVD) is a powerful tool in linear algebra that allows us to decompose a matrix into its constituent parts.

Remark

The singular value decomposition (SVD) is a powerful tool in linear algebra that allows us to decompose a matrix into its constituent parts.

In the case of the covariance matrix Σ , the SVD provides a way to express it in terms of its **singular values** (the diagonal entries of Λ , also known as its **eigenvalues**) and its **singular vectors** (the columns of Q , also known as its **eigenvectors**).

Remark

The singular value decomposition (SVD) is a powerful tool in linear algebra that allows us to decompose a matrix into its constituent parts.

In the case of the covariance matrix Σ , the SVD provides a way to express it in terms of its **singular values** (the diagonal entries of Λ , also known as its **eigenvalues**) and its **singular vectors** (the columns of Q , also known as its **eigenvectors**).

This decomposition is particularly useful for understanding the geometry of the multivariate normal distribution.

Notice that, since Σ is positive definite, the diagonal entries of

$$\Lambda = \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_p \end{bmatrix}$$

are all positive.

Notice that, since Σ is positive definite, the diagonal entries of

$$\Lambda = \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_p \end{bmatrix}$$

are all positive.

In particular, we can define a new matrix

$$D = \begin{bmatrix} \sqrt{\lambda_1} & 0 & \cdots & 0 \\ 0 & \sqrt{\lambda_2} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sqrt{\lambda_p} \end{bmatrix}$$

satisfying $D^2 = \Lambda$.

If we now take

$$\sigma = DQ^T$$

If we now take

$$\sigma = DQ^T$$

then

$$\Sigma = \sigma^T \sigma$$

If we now take

$$\sigma = DQ^T$$

then

$$\Sigma = \sigma^T \sigma$$

In particular, we can define the **standardization** of the random vector X as

$$Z = (X - \mu)\sigma^{-1} = (X - \mu)QD^{-1} = (X - \mu)Q\Lambda^{-1/2}$$

in which case $Z \sim \mathcal{N}_p(0, I_p)$.

Closure under Linear Transformations

Theorem

Let $X \sim \mathcal{N}_p(\mu, \Sigma)$ be a random vector with normal multivariate distribution. Then, for any matrix A of size $m \times p$ and any vector b of size m , the random vector

$$Y = AX + b$$

also has a normal multivariate distribution, specifically $Y \sim \mathcal{N}_m(A\mu + b, A\Sigma A^T)$.

Remark

This theorem states that if we apply a linear transformation to a random vector with a normal multivariate distribution, the resulting random vector will also have a normal multivariate distribution.

Remark

This theorem states that if we apply a linear transformation to a random vector with a normal multivariate distribution, the resulting random vector will also have a normal multivariate distribution.

The mean and covariance of the transformed vector are given by the formulas in the theorem.

Remark

This theorem states that if we apply a linear transformation to a random vector with a normal multivariate distribution, the resulting random vector will also have a normal multivariate distribution.

The mean and covariance of the transformed vector are given by the formulas in the theorem.

This property is particularly useful in applications such as regression analysis, where we often need to transform the data to fit a model while maintaining the normality of the residuals.

Remark

This theorem states that if we apply a linear transformation to a random vector with a normal multivariate distribution, the resulting random vector will also have a normal multivariate distribution.

The mean and covariance of the transformed vector are given by the formulas in the theorem.

This property is particularly useful in applications such as regression analysis, where we often need to transform the data to fit a model while maintaining the normality of the residuals.

In practice, this theorem is often used to derive the distribution of linear combinations of random variables, which is a common operation in multivariate statistics.

Example

Let $X \sim \mathcal{N}_2(\mu, \Sigma)$ be a random vector with mean $\mu = (1, 2)$ and covariance matrix

$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix}$$

Example

Let $X \sim \mathcal{N}_2(\mu, \Sigma)$ be a random vector with mean $\mu = (1, 2)$ and covariance matrix

$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix}$$

Now, consider the linear transformation

$$Y = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} X + \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Example

Let $X \sim \mathcal{N}_2(\mu, \Sigma)$ be a random vector with mean $\mu = (1, 2)$ and covariance matrix

$$\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix}$$

Now, consider the linear transformation

$$Y = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} X + \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

In other words,

$$Y_1 = 2X_1 + X_2$$

$$Y_2 = 3X_2 + 1$$

Example

Using the closure under linear transformations theorem, we can find the mean and covariance of Y :

$$\text{Mean: } A\mu + b = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 4 \\ 7 \end{bmatrix}$$

$$\text{Covariance: } A\Sigma A^T = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 1 & 3 \end{bmatrix} = \begin{bmatrix} 8 & 9 \\ 9 & 18 \end{bmatrix}$$

Example

Using the closure under linear transformations theorem, we can find the mean and covariance of Y :

$$\text{Mean: } A\mu + b = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 4 \\ 7 \end{bmatrix}$$

$$\text{Covariance: } A\Sigma A^T = \begin{bmatrix} 2 & 1 \\ 0 & 3 \end{bmatrix} \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 1 & 3 \end{bmatrix} = \begin{bmatrix} 8 & 9 \\ 9 & 18 \end{bmatrix}$$

Thus,

$$Y \sim \mathcal{N}_2 \left(\begin{bmatrix} 4 \\ 7 \end{bmatrix}, \begin{bmatrix} 8 & 9 \\ 9 & 18 \end{bmatrix} \right)$$

Remark

Notice that since projection into a subspace is a linear transformation, the closure under linear transformations theorem implies that any subset of the components of a normal multivariate distribution also has a normal multivariate distribution.

Remark

Notice that since projection into a subspace is a linear transformation, the closure under linear transformations theorem implies that any subset of the components of a normal multivariate distribution also has a normal multivariate distribution.

This is a very useful property, as it allows us to analyze subsets of variables while still maintaining the normality assumption.

Multivariate Normality Testing

There are several statistical tests to check whether a given dataset follows a multivariate normal distribution. These tests are essential in multivariate statistics, as many methods assume normality. Here are some of the most commonly used tests, together with their implementation in R:

1. **Mardia's Test:** Based on multivariate skewness and kurtosis. For a multivariate normal distribution, skewness should be close to 0 and kurtosis close to $p(p+2)$, where p is the number of variables. Implemented in R via `MVN::mvn(your_data, mvn_test = "mardia")`.
2. **Henze-Zirkler (HZ) Test:** Uses the empirical characteristic function to measure deviations from normality. It has good global power and is suitable for moderate dimensions. Implemented in R via `MVN::mvn(your_data, mvn_test = "hz")`.

3. **Royston's Test:** A multivariate extension of the Shapiro-Wilk test, based on the product of univariate Shapiro-Wilk statistics applied to principal components. Easy to interpret, though power decreases in high dimensions. Implemented in R via `MVN::mvn(your_data, mvn_test = "royston")`.
4. **Doornik-Hansen Test:** Combines skewness and kurtosis into a single omnibus test, providing a robust alternative to Mardia. Implemented in R via `MVN::mvn(your_data, mvn_test = "dh")`.
5. **Graphical Methods (Chi-square Q-Q plots):** Plot the Mahalanobis distances of the data against the quantiles of a chi-square distribution. Deviations from a straight line indicate departures from multivariate normality. Implemented in R via `multivariate_diagnostic_plot(data = your_data, type = "qq")`.